Linear Regression Models P8111

Lecture 21

Jeff Goldsmith April 12, 2016



Today's Lecture

- LDA review
- Model interpretation
- Model estimation
- Example (pig data!)

Recall the setting

- We observe data y_{ij} , x_{ij} for subjects i = 1, ..., I at visits $j = 1, ..., J_i$
- Overall, we pose the model

$$y = X\beta + \epsilon$$

where $Var(\epsilon) = \sigma^2 V$ and

$$V = \left[\begin{array}{cccc} V_1 & 0 & \dots & 0 \\ 0 & V_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \\ 0 & 0 & & V_I \end{array} \right]$$

General ideas

- We discussed two approaches:
 - Random (mixed) effects models introduce random subject effects and assume uncorrelated errors
 - Marginal models include only fixed effects and assume correlated subject-level errors
- In some cases these introduce equivalent correlation structures; an example is that a random intercept model induces uniform within-subject correlations
- Marginal models separate the mean structure and correlation; random effect models induce correlation by introducing random quantities into the mean

Some comparisons

- Marginal models may be less immediately obvious or interpretable
- However, marginal models can be more robust to misspecification through robust SE estimates

Marginal model

The marginal model formulation is

$$y = X\beta + \epsilon$$

where

- $lacksquare \epsilon \sim \mathcal{N}\left[0,
 u^2 V\right]$
- This approach focuses on the *marginal* distribution of *y*, rather than on a subject-level *conditional* distribution
- Coefficients have a marginal interpretation compare subjects based only on covariate values
- Interpretation is analogous to a cross-sectional model

Remember GLS

Given the model

$$y = X\beta + \epsilon$$

where $\epsilon \sim N(0, \sigma^2 V)$ with V known, we are essentially assuming

$$y \sim N(X\beta, \sigma^2 V)$$

Using MLE, we find that $\hat{\beta}_{GLS} = (X^T V^{-1} X)^{-1} X^T V^{-1}$

Estimation – marginal model

- If we can use MLE when *V* is known, maybe we can use MLE to estimate *V* as well
- Our log likelihood function is

$$l(\beta, \sigma^2, V; y, X) = -\frac{1}{2} \left[n \log(\sigma^2) + \log(|V|) + \frac{1}{\sigma^2} (y - X\beta)^T V^{-1} (y - X\beta) \right]$$

• Using profile likelihood, we find that for any V_0

$$\hat{\boldsymbol{\beta}}(V_0) = (\boldsymbol{X}^T V_0^{-1} \boldsymbol{X})^{-1} \boldsymbol{X}^T V_0^{-1}$$

Estimation – marginal model

- Estimation of V and σ is done through restricted maximum likelihood
 - Standard MLE produces biased variance estimates; REML adjusts for the number of fixed effects components that are estimated
- Often V is structured parametrically to ease estimation and computation
- We won't worry about how this is done

A random intercept model with one covariate is given by

$$y_{ij} = \beta_0 + b_i + \beta_1 x_{ij} + \epsilon_{ij}$$

where

- $b_i \sim N [0, \tau^2]$
- \bullet $\epsilon_{ij} \sim N \left[0, \nu^2\right]$

More compactly, we write

$$y = X\beta + Zb + \epsilon$$

where

- $lackbox{b} \sim \mathrm{N}\left[0, au^2 I_I\right]$
- $\epsilon \sim N \left[0, \nu^2 I_n\right]$

In the model

$$y = X\beta + Zb + \epsilon$$

why might we assume b are random rather than estimating them as fixed effects?

In the model

$$y = X\beta + Zb + \epsilon$$

why might we assume *b* are fixed rather than considering them to be random?

Random intercept model interpretation

Random effect models have a conditional interpretation.

- Mean conditions on subject effects: $E(y_{ij}|x_{ij}, \beta, b_i)$
- To derive a marginal (averaged across subjects) mean, one can use iterated expectations

■ (For an identity link, the coefficients also have a marginal interpretation – this is not true for generalized RE models)

Estimation – random effect model

Still done using MLE, but now we include random effects; our model is

$$y = X\beta + Zb + \epsilon$$

where

- lacksquare $b \sim N \left[0, \tau^2 I_I\right]$
- \bullet $\epsilon \sim N \left[0, \nu^2 I_n\right]$

Estimation – random effect model

Estimation – random effect model

Estimation – BLUPs

Our estimate for fixed and random effects are

$$\begin{bmatrix} \hat{\beta} \\ \hat{b} \end{bmatrix} = \left(C^T C + \frac{\nu^2}{\tau^2} R \right)^{-1} C^T y$$

- These are referred to as "BLUPs" (the "P" is for "predictions")
- The variances ν^2 and τ^2 are estimated via REML

BLUPs

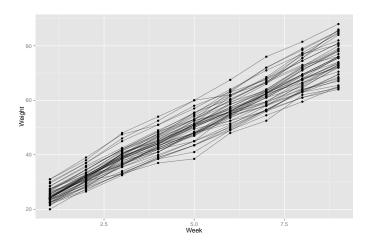
BLUPs are a lot like BLUEs

- They are the best linear unbiased predictors when one has both fixed and random effects
- We derived them using Normal distributions, but even without distributional assumptions these are BLUP
- The Normal distributions help with the assumptions are satisfied, in that one can get distribution-based inference

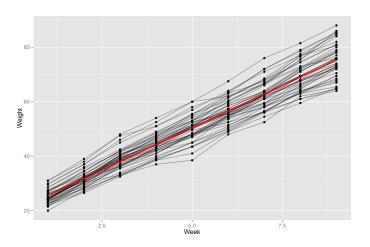
Penalized spline regression

Note that BLUPs look an awful lot like penalized regression estimates

- The "mixed model framework" is commonly used for penalized spline estimation
- The "tuning parameter" is a ratio of variances, and can be estimated via REML (rather than CV)
- This approach provides a method for inference

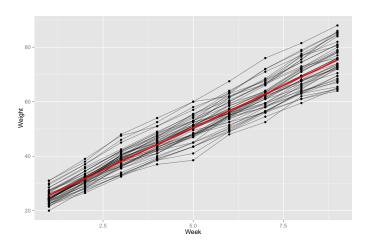


OLS fit for pig data



OLS code

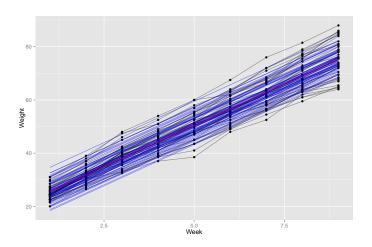
Marginal model fit for pig data



Marginal model code

```
> library(gee)
> marg.mod = gee(weight ~ num.weeks, id = id.num, corstr = "exchangeable",
                data = pig.weights)
> summary (marg.mod)
Model .
Link.
                           Identity
Variance to Mean Relation: Gaussian
 Correlation Structure: Exchangeable
Summary of Residuals:
       Min
                    10 Median 30
                                                      Max
-11.9050926 -2.5347801 -0.1951968 2.5949074 13.1751157
Coefficients.
            Estimate Naive S.E. Naive z Robust S.E. Robust z
(Intercept) 19.355613 0.5983680 32.34734 0.39963854 48.43280
num.weeks 6.209896 0.0393321 157.88366 0.09107443 68.18485
Working Correlation
                                       [,4] ...
 [1,] 1.0000000 0.7690313 0.7690313 0.7690313 ...
 [2,] 0.7690313 1.0000000 0.7690313 0.7690313 ...
 [3,] 0.7690313 0.7690313 1.0000000 0.7690313 ...
 [4,] 0.7690313 0.7690313 0.7690313 1.0000000 ...
```

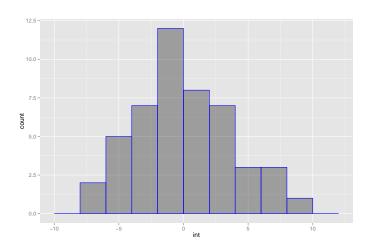
Mixed effect model fit for pig data



Mixed effect model code

```
> library(lme4)
> ranef.mod = lmer(weight ~ (1 | id.num) + num.weeks, data = pig.weights)
> summary(ranef.mod)
Linear mixed model fit by REML
Formula: weight ~ (1 | id.num) + num.weeks
  Data: pig.weights
 AIC BIC logLik deviance REMLdev
 2042 2058 -1017 2030 2034
Random effects:
Groups Name Variance Std.Dev.
id.num (Intercept) 15.1418 3.8913
Residual
                    4.3947 2.0964
Number of obs: 432, groups: id.num, 48
Fixed effects.
           Estimate Std. Error t value
(Intercept) 19.35561 0.60311 32.09
num.weeks 6.20990 0.03906 158.97
> (15.1418) / (15.1418 + 4.3947)
[1] 0.7750518
```

Histogram of estimated random intercepts



Today's big ideas

- Estimation in LDA
- Example + code

Potential reading on mixed effects models – Semiparametric Regression (Ruppert, Wand, Carroll) Ch 4